




Machine Learning - part 2

Deep Learning



GGE 6505/GGE5405 Introduction to
Big Data & Data Science



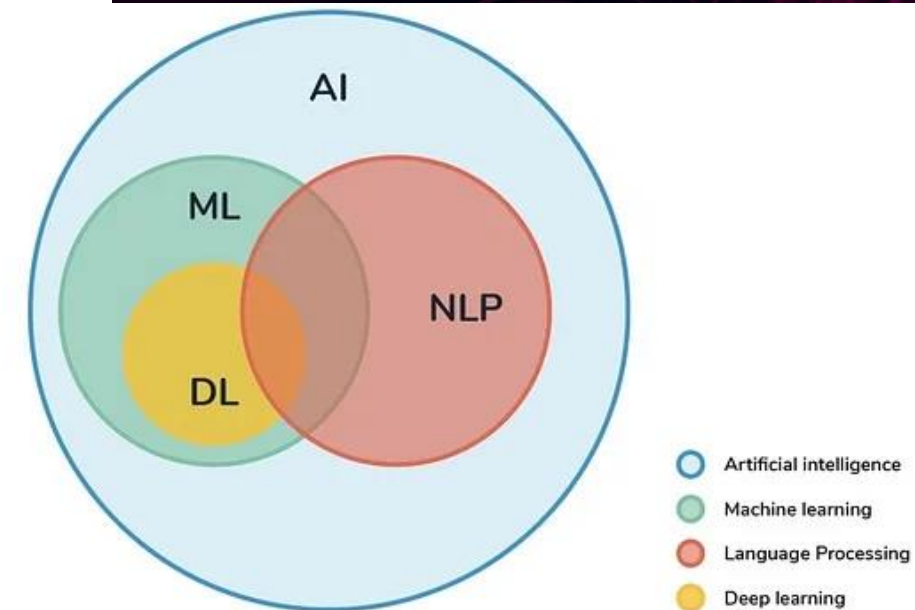
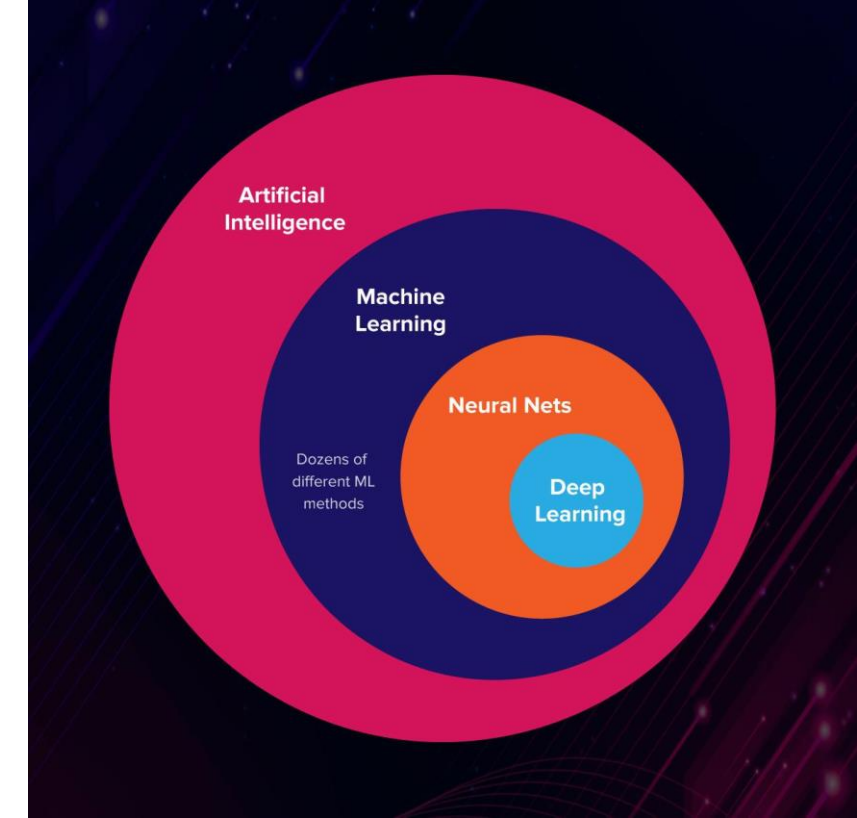
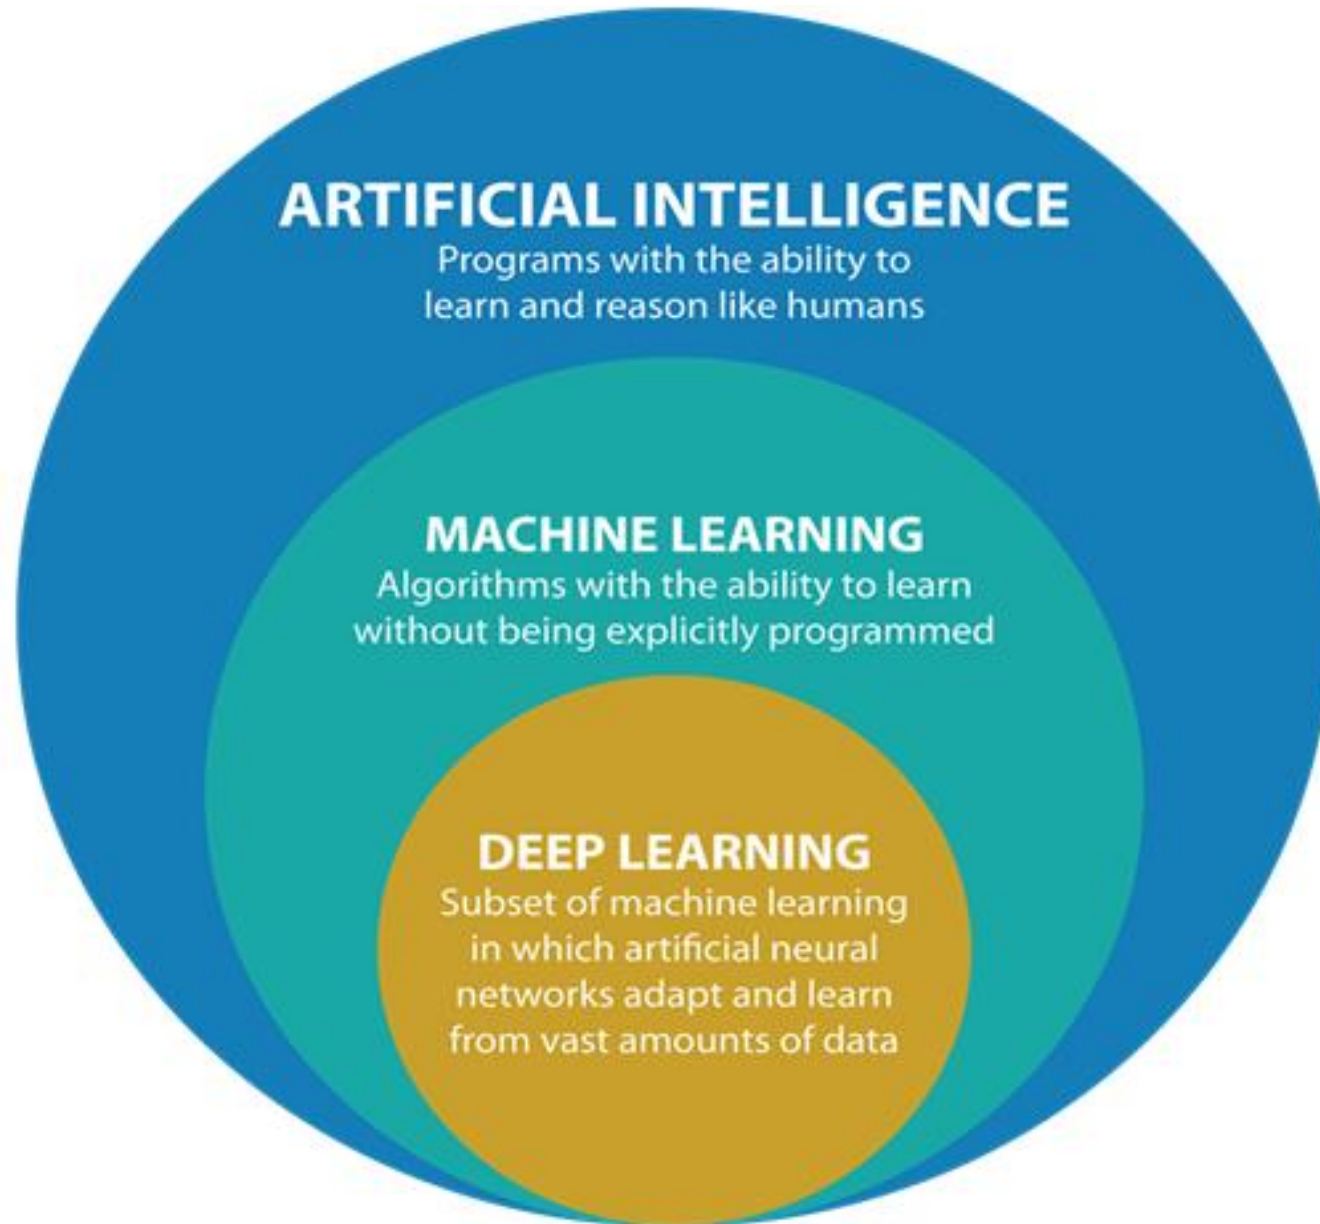
Deep Learning Textbook

- [Deep Learning \(deeplearningbook.org\)](http://deeplearningbook.org)

Goodfellow-et-al-2016, Deep Learning, Ian Goodfellow and Yoshua Bengio and Aaron Courville, MIT Press, 2016



Machine learning vs AI vs Deep learning



Machine learning vs AI vs Deep learning

AI: Artificial intelligence is a field of computer science that makes a computer system that can mimic human intelligence

Machine learning: Machine learning enables a computer system to make predictions or take some decisions using historical data without being explicitly programmed

Deep Learning/ Neural Networks: The biggest advantage of Deep Learning is that we do not need to manually extract features from the image. The network learns to extract features while training. You just feed the image to the network (pixel values).

AI VS. MACHINE LEARNING VS. DEEP LEARNING


Artificial Intelligence: a program that can sense, reason, act and adapt.

Machine Learning: algorithms whose performance improve as they are exposed to more data over time.

Deep Learning: subset of machine learning in which multilayered neural networks learn from vast amounts of data.

Exciting progress:



- Face recognition
- Image classification
- Speech recognition
- Text-to-speech generation
- Handwriting transcription
- Machine translation
- Medical diagnosis
- Cars: drivable area, lane keeping
- Digital assistants
- Ads, search, social recommendations
- Game playing with deep RL




Machine Learning




Deep Learning



Neuron: Biological Inspiration for Computation



- **Neuron:** computational building block for the brain

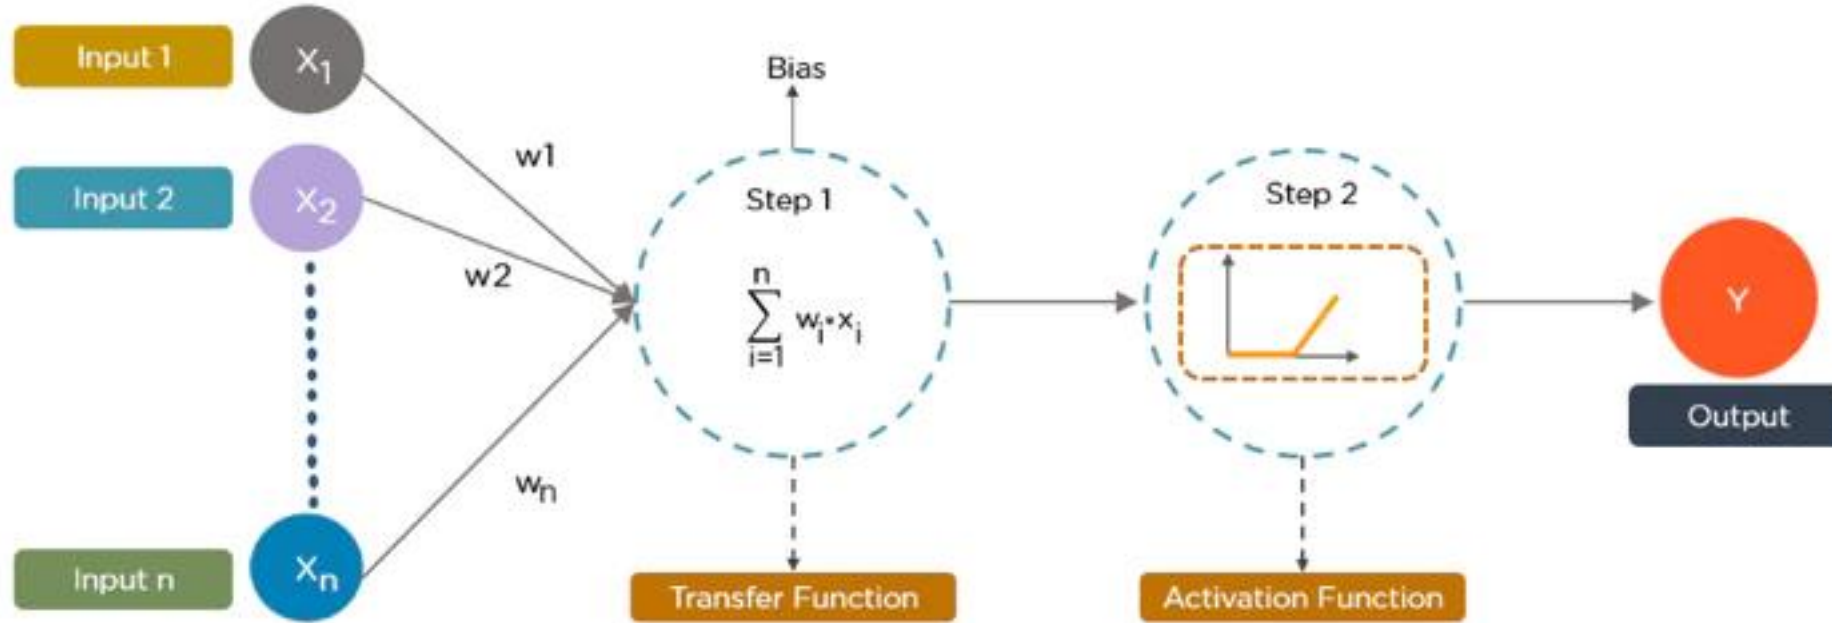


- **(Artificial) Neuron:** computational building block for the "neural network"

Key Difference:


- **Parameters:** Human brains have $\sim 10,000,000$ times synapses than artificial neural networks.
- **Topology:** Human brains have no "layers". **Async:** The human brain works asynchronously, ANNs work synchronously.
- **Learning algorithm:** ANNs use gradient descent for learning. We don't know what human brains use
- **Power consumption:** Biological neural networks use very little power compared to artificial networks
- **Stages:** Biological networks usually never stop learning. ANNs first train then test.

Neural Network




[Top 10 Deep Learning Algorithms You Should Know in 2022 \(simplilearn.com\)](https://simplilearn.com)

Neural Network can be considered a class of algorithms, mechanisms, or models of Machine Learning that take their inspiration from a human brain since they can have simple to highly complex 'neural' networks. Having these layers helps a machine perform computations that would not be possible using just the standard ML algorithms.



Forward Propagation



Key concepts:

Activation Function

Activation functions in general are used to convert linear outputs of a neuron into [nonlinear outputs](#), ensuring that a neural network can learn nonlinear behavior.

```
import math


def sigmoid(x):
    return 1 / (1 + math.exp(-x))
```

```
def tanh(x):
    return (math.exp(x) - math.exp(-x)) / (math.exp(x) + math.exp(-x))
```

```
def relu(x):
    return max(0, x)
```


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$



ReLU

$$\max(0, x)$$



Sigmoid function

Advantages:


- It introduces non linearity into dataset
- It can be differentiated which will be resulted into backpropagation in neural network.
- Normalize the data by confining it in a range of [0,1].

Disadvantages:

- Resulted into Vanishing Gradient problem.
- Shows Non-Zero centric behavior.
- Due to presence of exponential function it becomes computationally too much Expensive.

Mathematical Representation

- $f(x) = 1 \div (1 + \exp(-x))$



Hypertangent function(tanh)

Advantages:


- Zero centric Function
- finds new weight and biases with every back propagation because it's derivative lies in a range [0,1]
- Better than sigmoid function.

Disadvantages:

- Vanishing Gradient problem.

Mathematical Representation

$$f(x) = (\exp(x) - \exp(-x)) \div (\exp(x) + \exp(-x))$$



ReLU (Rectified Linear Activation)

ReLU overcomes the Vanishing Gradient problem which is found in Sigmoid and tanh function


Advantages:

- It doesn't activate all the neuron at the same time


Disadvantages:

- If the value of training examples will be negative ($x < 0$) it will not be able to find any gradient hence resulted in dying Perceptron problem.


$$\sigma(x) = \begin{cases} \max(0, x) & , x \geq 0 \\ 0 & , x < 0 \end{cases}$$



Multiple Hidden Layers



ReLU
 $\max(0, x)$




Loss Functions




Regression

What is the temperature going to be tomorrow?



Classification

Will it be Cold or Hot tomorrow?



- Loss function quantifies gap between prediction and ground truth
- For regression:
 - Mean Squared Error (MSE)
- For classification:
 - Cross Entropy Loss

Mean Squared Error

$$MSE = \frac{1}{N} \sum (t_i - s_i)^2$$

Prediction (points to s_i)

Ground Truth (points to t_i)

Cross Entropy Loss


$$CE = - \sum_i^C t_i \log(s_i)$$

Classes (points to C)

Prediction (points to s_i)



Ground Truth {0,1} (points to t_i)

Backpropagation



Actual Value of target : 13
Error: Predicted – Actual = -4

Backpropagation



Task: Update the **weights** and **biases** to decrease **loss function**


Subtasks:

1. Forward pass to compute network output and “error”
2. Backward pass to compute gradients
3. A fraction of the weight’s gradient is subtracted from the weight.

↑
Learning Rate

Numerical Method: **Automatic Differentiation**

Simple Neural Network




● Input Layer

● Hidden Layer

● Output Layer


Deep Learning Neural Network



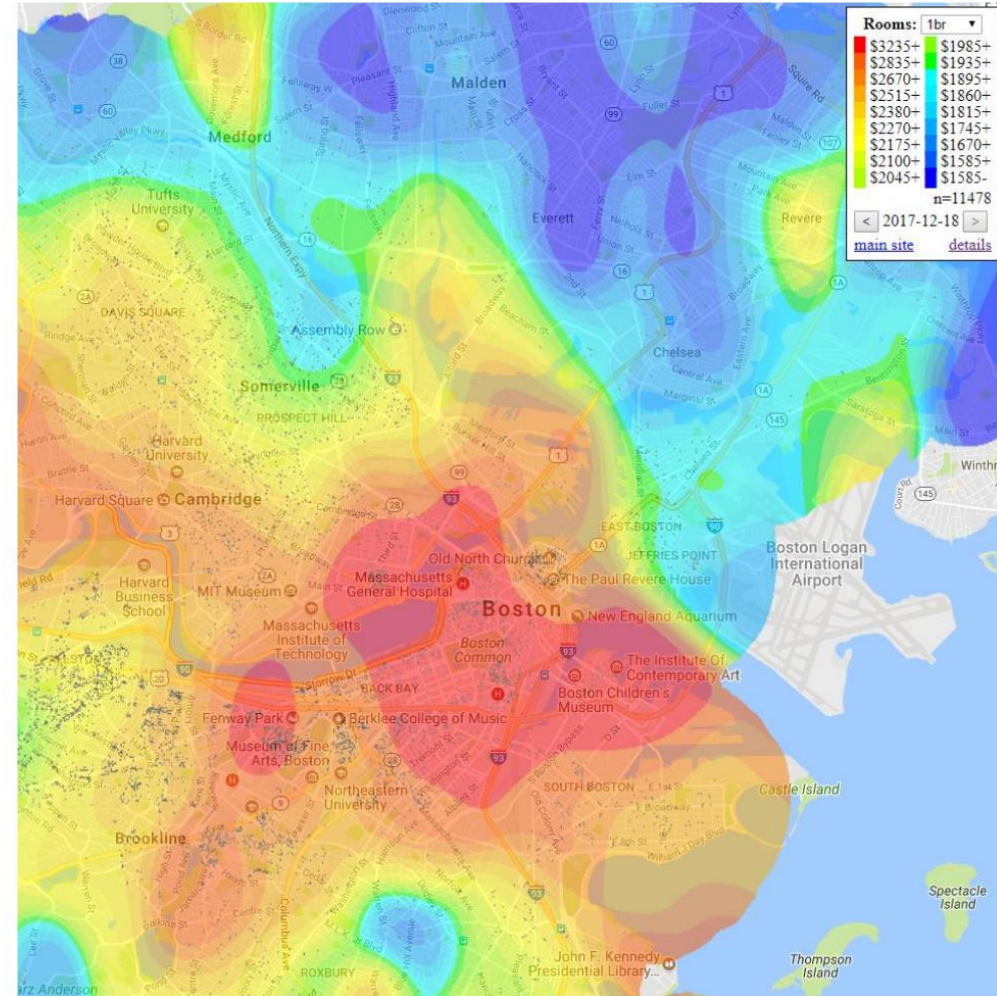
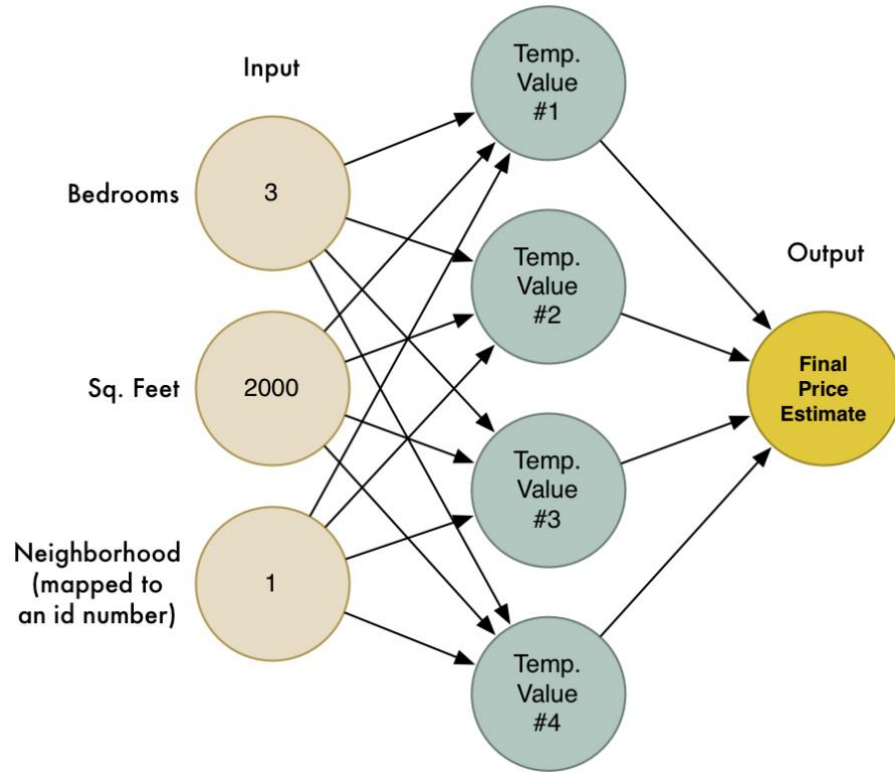
Universality: For any arbitrary function $f(x)$, there exists a neural network that closely approximate it for any input x

Neural Network

- Housing price prediction



Special Purpose Intelligence: Estimating Apartment Cost



A Neural Network Playground (tensorflow.org)



Epoch
000,000

Learning rate
0.03

Activation
Tanh

Regularization
None

Regularization rate
0

Problem type
Classification

DATA

Which dataset do you want to use?



Ratio of training to test data: 50%

Noise: 0

Batch size: 10


REGENERATE

FEATURES

Which properties do you want to feed in?

- X_1
- X_2
- X_1^2
- X_2^2
- $X_1 X_2$
- $\sin(X_1)$
- $\sin(X_2)$

2 HIDDEN LAYERS




This is the output from one neuron. Hover to see it larger.


The outputs are mixed with varying weights, shown by the thickness of the lines.

OUTPUT

Test loss 0.489
Training loss 0.498



Colors shows data, neuron and weight values.




☐ Show test data ☐ Discretize output



TensorFlow in One Slide

- **What is it:** Deep Learning Library (*and more*)
 - **Facts:** Open Source, Python, Google
- **Community:**
 - 117,000+ GitHub stars
 - TensorFlow.org: Blogs, Documentation, DevSummit, YouTube talks
- **Ecosystem:**
 - **Keras:** high-level API
 - **TensorFlow.js:** in the browser
 - **TensorFlow Lite:** on the phone
 - **Colaboratory:** in the cloud
 - **TPU:** optimized hardware
 - **TensorBoard:** visualization
 - **TensorFlow Hub:** graph modules
- **Extras:**
 - Swift for TensorFlow
 - TensorFlow Serving
 - TensorFlow Extended (TFX)
 - TensorFlow Probability
 - Tensor2Tensor
- **Alternatives:** PyTorch, MXNet, CNTK



- 1

```
# import tensorflow and keras (tf.keras not "vanilla" Keras)
import tensorflow as tf
from tensorflow import keras
```
- 2

```
# get data
(train_images, train_labels), (test_images, test_labels) = \
keras.datasets.mnist.load_data()
```
- 3

```
# setup model
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28, 28)),
    keras.layers.Dense(128, activation=tf.nn.relu),
    keras.layers.Dense(10, activation=tf.nn.softmax)
])

model.compile(optimizer=tf.train.AdamOptimizer(),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```
- 4

```
# train model
model.fit(train_images, train_labels, epochs=5)
```
- 5

```
# evaluate
test_loss, test_acc = model.evaluate(test_images, test_labels)
print('test accuracy:', test_acc)
```
- 6

```
# make predictions
predictions = model.predict(test_images)
```


Deep learning models:

Some of the top Deep Learning Algorithms and Models include the following:

Convolutional Neural Network (CNN)

Radial Basis Function Networks (RBFNs)

Multilayer Perceptrons (MLPs)

Self Organizing Maps (SOMs)

Deep Belief Networks (DBNs)

Restricted Boltzmann Machines (RBMs)

Autoencoders

Convolutional Neural Networks (CNNs)

Long Short-Term Memory Networks (LSTMs)

Recurrent Neural Networks (RNNs) tx

Generative Adversarial Networks (GANs)

Convolutional Neural Networks: Image Classification






image classification

82% cat
15% dog
2% hat
1% mug






0	0	0	0	0	...	0
0	87	240	210	24	...	0
0	13	0	101	195	...	0
0	35	167	99	210	...	0
0	145	230	240	201	...	140
...
0	0	0	0	0	...	0

28 by 28 grid




$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ \dots \\ 0 \\ 0 \\ 87 \\ 240 \\ 210 \\ 24 \\ 0 \\ \dots \\ 0 \end{bmatrix}$



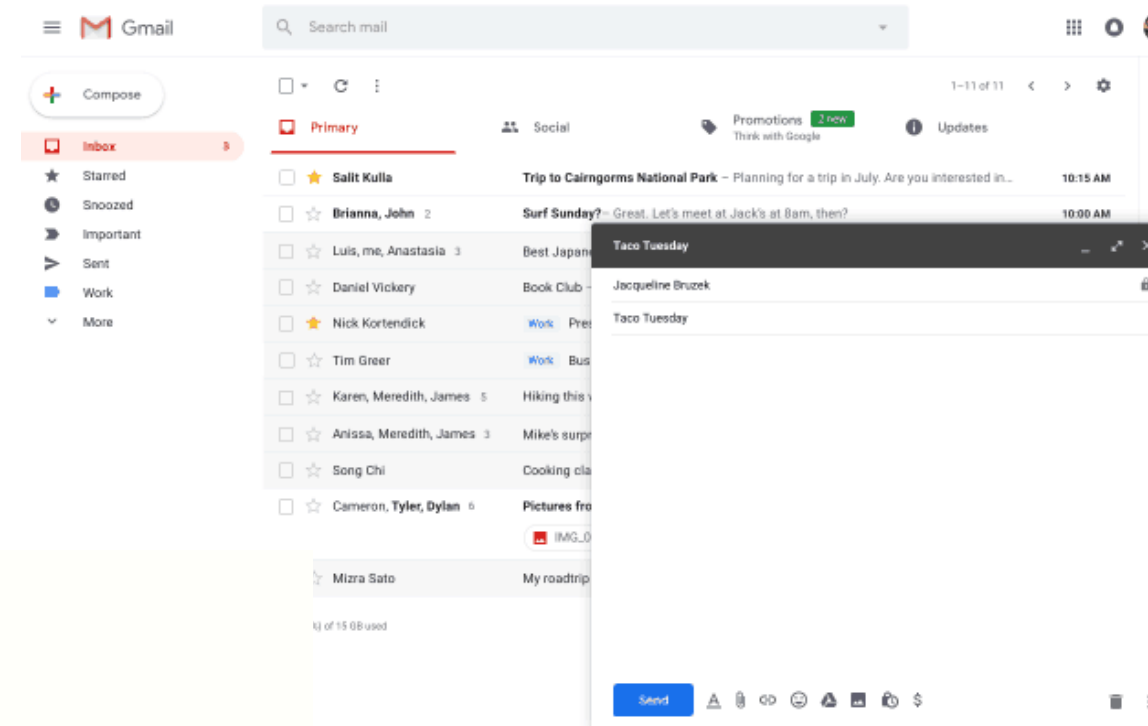
Recurrent Neural Network (RNN)

A recurrent neural network (RNN) is a **special type of an artificial neural network adapted to work for time series data or data that involves sequences**




- Applications

- Sequence Data
- Text
- Speech
- Audio
- Video
- Generation



How do we get a feed-forward neural network to be able to use previous information to effect later ones?
What if we add a loop in the neural network that can pass prior information forward?

What time is it?



References and materials to read

- [3Blue1Brown – YouTube](#)
- [codebasics – YouTube](#)
- [Lex Fridman - YouTube](#)